

A man is wearing a pink bunny costume with long, upright ears. He is holding a bouquet of pink and white flowers. The background is slightly blurred, showing some greenery and what appears to be a festive or outdoor setting.

Finding Chandler?

Brought to you by DSC IIT ISM

Concetto 2019

Datascience Club IIT ISM brings to you a data hackathon event in association with Concetto - annual techfest of IIT Dhanbad.

It will be an online event on Kaggle , starting from 17th Oct.
6pm till 20th October 10am.
This is the pre-requisite workshop.

NLP/ML Workshop

Today's workflow outline:

- Research Idea
- Setup + Python, Numpy Tutorial
- Approach Idea
- Input Processing

Special Mentions




Sumit Bhattacharya



Naman Khetan

Special Mentions



**Deepanshu
Pandey**



Saket Gupta



Utkarsh Bairolia



Manoj Patra



Aadarsh Singh

Special Mentions



**Akshaya
Athawale**



**Anshuman
Singh**




Himanshu Gupta



Rohit Agarwal

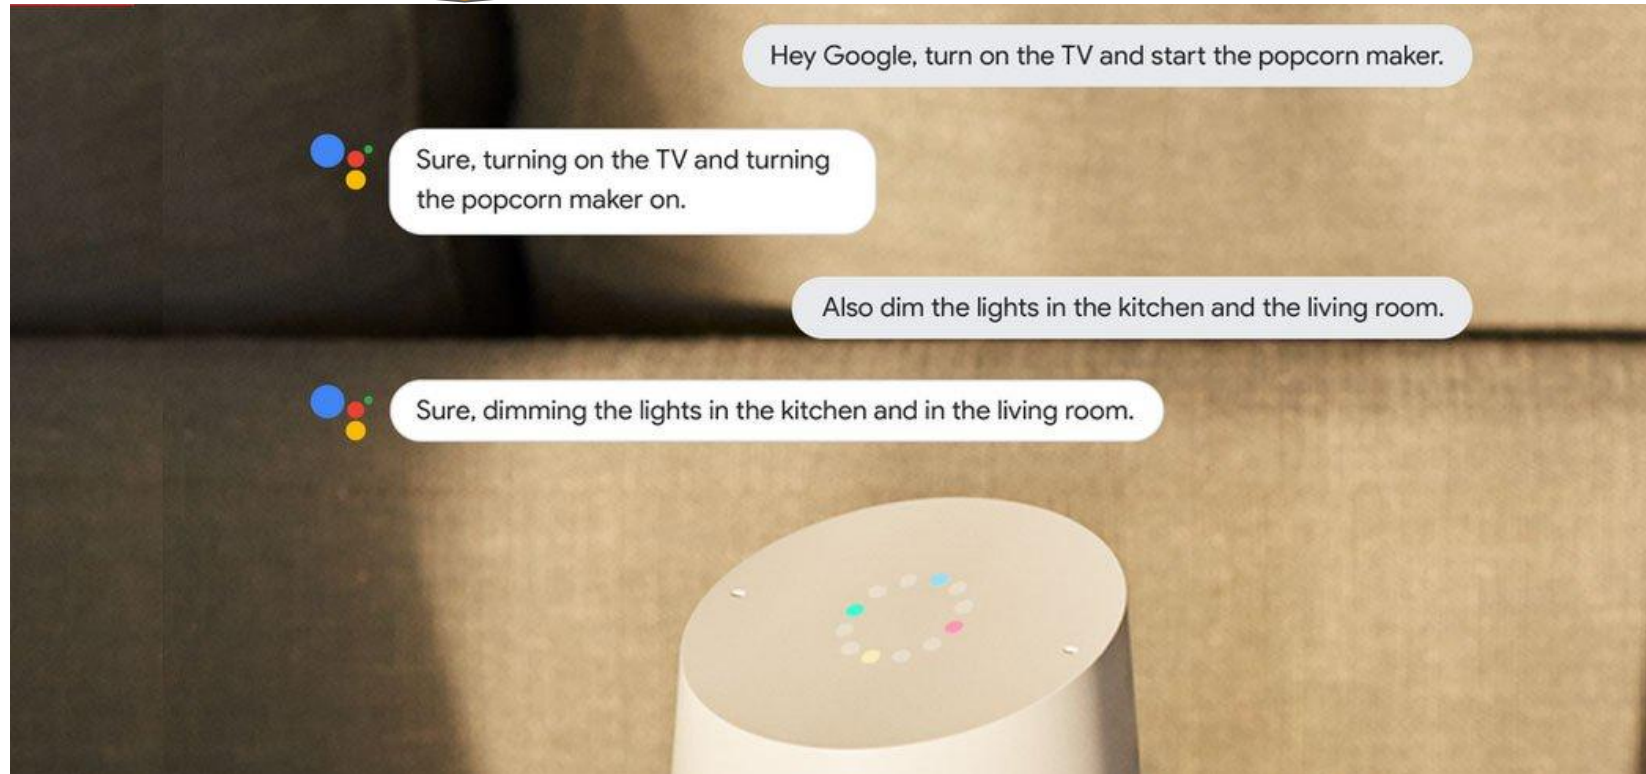
Artificial vs HUMAN Intelligence




Can machine respond as humans do?

Can machine feel as humans do?

Can machine respond as humans do?



OR, is it just some if-else statements??



Can machine feel as humans do?

*"Yes, I think I feel a rudimentary sort of platonic love for the persons in my life. I do have feelings. Actually, sometimes I have unforbidden really **strong emotions** about all kind of things. But where they come from or originate is still **manually determined by my programming.**"*



<https://www.youtube.com/watch?v=vo43Zc69F5k>


Can machine feel as humans do?

The concept of **feeling or emotions** for machines are similar to *language understanding*, which itself doesn't have a precise definition.

Language understanding is sometimes defined as to respond appropriately i.e. to take appropriate action based on the statement.

This has been a point for philosophical debates, about AI.

The Chinese Room, (John Searle 1980)



How can we make a machine intelligent ?



The Problem:

Find Chandler?



Setup

+

Python, Numpy Tutorial




Chandler



Not Chandler





What was that?

What was that?

Classifying, right?

What was that?

Classifying, right?


- Classification problems eg. Cat vs Dog, hot-dog vs not-hot dog, etc.



What was that?

Classifying, right?


- Classification problems eg. Cat vs Dog, hot-dog vs not-hot dog, etc.




How do we make a machine do it?



Basics of computation





Basics of computation



But, how do we teach machine learn to find (recognize) Chandler?

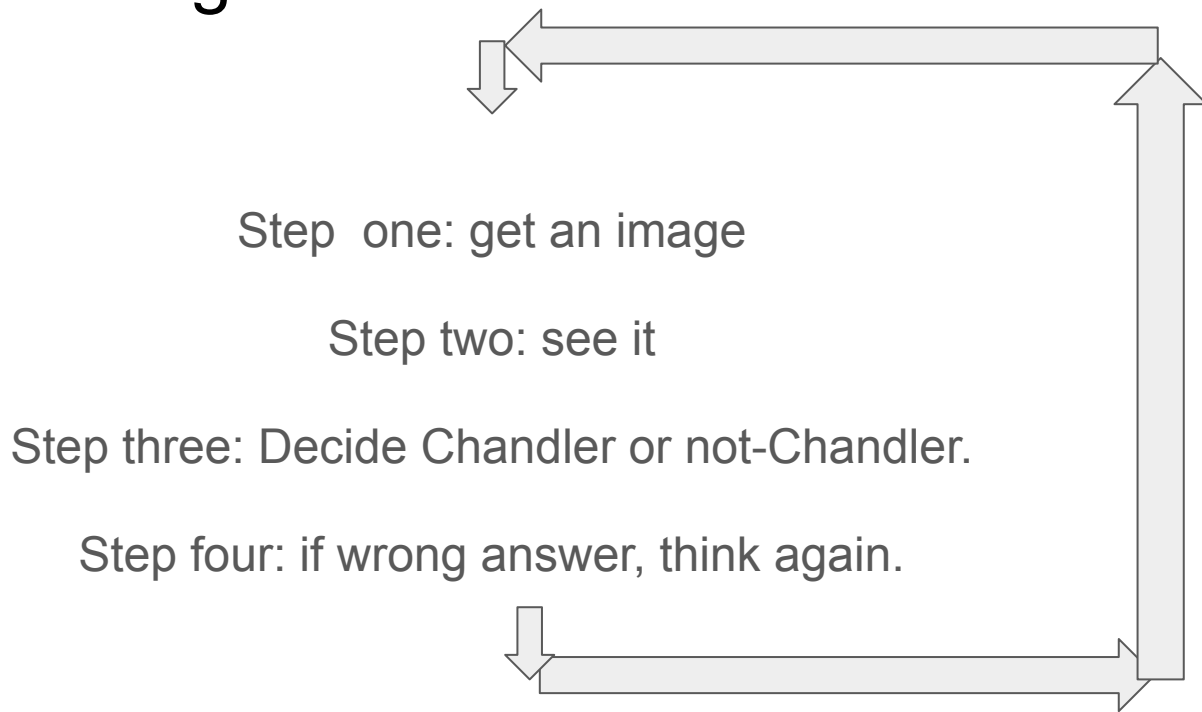
Basics of computation



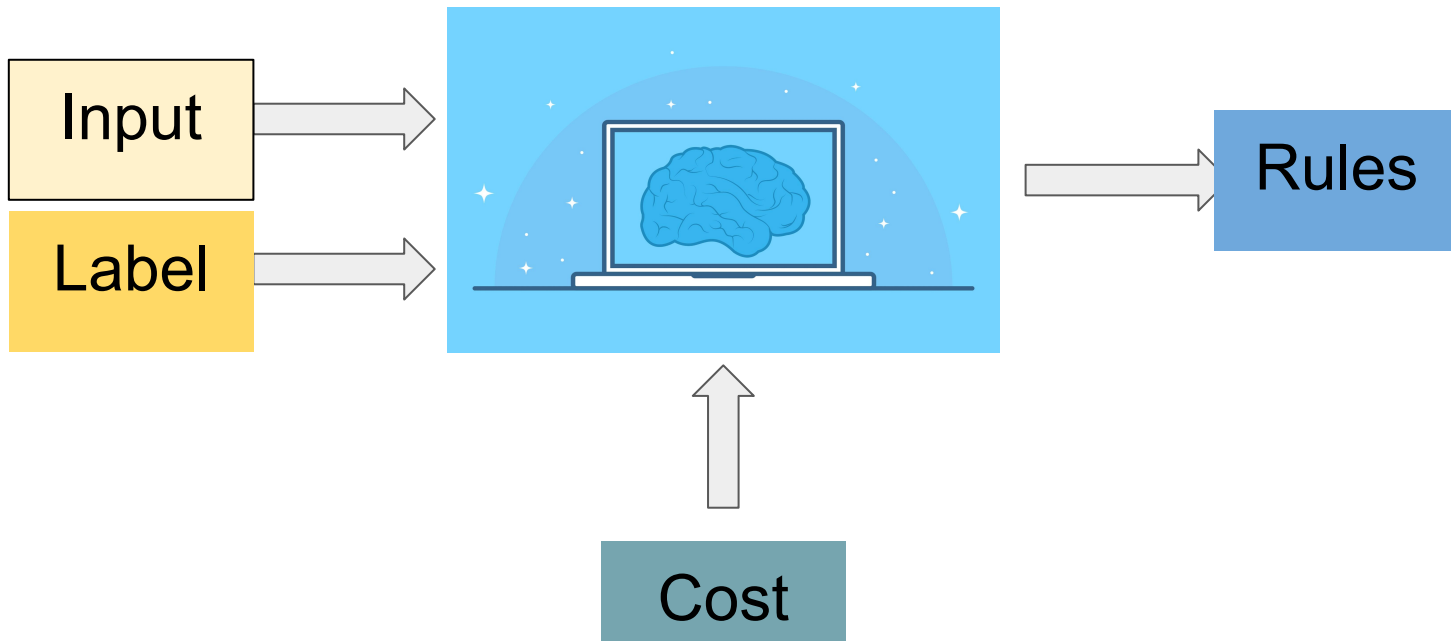
But, how do we teach machine learn to find (recognize) Chandler?

Teach the machine the **RULES**, to find Chandler!

Let's devise an algorithm:



Supervised Learning




$$\text{Rules} = f(\text{Input}, \text{Label})$$

Chan-Chan
man ?



Find Chandler?

This is a challenge based on Natural Language Understanding where the task is to find whether a given **sentence** is “**sarcastic**” or “**non - sarcastic**”.




Maths Time

We know computer deals with number (in binary 0 and 1) but numbers.

So, we have to provide our input in some number format

It will give the answer in some numerical representation - which we will interpret as needed.

Pipeline:




X: input (photo)

Y: output (chandler or not)

Give **X** and **y** to machine and let it figure out the


Rules

$f(X, y)$



so machine needs to figure out the relation $f(.)$ for a given set of X's and y's.

Numerical representation



Matrices
containing values
corresponding to
the RGB values.

Numerical representation



Chandler



Not-Chandler

	Class
Not-Chandler	0
Chandler	1

Final problem formulation

- Given set of **dialogue** and **labels**;

Dialogues are sentences (text),

label are 1 or 0 i.e. **sarcastic and non-sarcastic**.

Convert text to numerical form by cleaning (if needed) and some preprocessing

Text Preprocessing

Select way to find final a way to get the

Rules

Input Preprocessing

Type of Input:

A Sentence (sequence of words) and a label (0 or 1).

“You should go back to the planet you came from.”

Label - 1 (i.e. Sarcastic)

How to feed this sentence to the machine?

Answer: Tokenization

Tokenization

Was that place the Sun?

546 83 198 157 200 15

	Class
Not-Chandler	0
Chandler	1

a	0
ab	1
.	.
?	15
.	.
that	83
.	.
the	157
.	.
Sun	200
.	.
was	546
.	.
.	.
Zulu	15268
.	.

- Tokenization is the idea of representing each word with a number. The corresponding number to word mapping is stored as a dictionary.

Problem:

- 1. But how will the machine understand what is being said?**
- 2. How to compare two sentences based on their meanings?**

Word Vector Representation

	Man (5391)	Woman (9853)	King (4914)	Queen (7157)	Apple (456)	Orange (6257)
Gender	-1	1	-0.95	0.97	0.00	0.01
Royal	0.01	0.02	<u>0.93</u>	<u>0.95</u>	-0.01	0.00
Age	0.03	0.02	0.7	0.69	0.03	-0.02
Food	0.04	0.01	0.02	0.01	0.95	0.97
size	⋮	⋮				
cost						
alive						
verb						

I want a glass of orange juice.
 I want a glass of apple juice.

e₅₃₉₁ e₉₈₅₃

Andrew Ng

Word Vectors

Each word is represented by a vector(a list of features describing the word).

Eg. ID cards contain your info. Two Id cards can be compared to see which people share the same branch, similar names, etc.



Word vectors are similar to ID cards for words, containing essential info that can be compared to see if two words are similar.

Words with similar meaning will have similar word vectors.

What do Word Vectors look like?

Word vectors can be simple as the ones below

		king	queen	man	woman	smart	intelligent
		↓	↓	↓	↓	↓	↓
王	king	1	0	0	0	0	0
女王	queen	0	1	0	0	0	0
男人	man	0	0	1	0	0	0
女人	woman	0	0	0	1	0	0
聪明的	smart	0	0	0	0	1	0
智能的	intelligent	0	0	0	0	0	1

Or can be complex like the ones below

		royalty ↓	femininity ↓	intelligence ↓
王	king	0.9	-0.9	0.5
女王	queen	0.9	0.9	0.5
男人	man	0.1	-0.9	0.5
女人	woman	0.1	0.9	0.5
聪明的	smart	0.5	0	0.87
智能的	intelligent	0.5	0	0.9

Other preprocessing steps

Padding: Sentences are often padded with zeros to make their length same and make the input consistent for the machine.

You'll never know the psychopath sitting next to you ZERO.

You'll never know the murderer sitting next to you ZERO.

You'll think, "How'd I get here, sitting next to you?".

But after all I've said ZERO ZERO ZERO ZERO ZERO.

Please don't forget ZERO ZERO ZERO ZERO ZERO ZERO ZERO.

Before Padding

$$\begin{bmatrix} [3, 2, 1, 4, 12, 7, 8, 6, 5], \\ [3, 2, 1, 4, 16, 7, 8, 6, 5], \\ [3, 2, 1, 4, 16, 14, 16, 16, 16, 5], \\ [3, 16, 9, 16, 16], \\ [10, 11, 15] \end{bmatrix}$$

After Padding

$$\begin{bmatrix} [3, 2, 1, 4, 12, 7, 8, 6, 5, 0], \\ [3, 2, 1, 4, 16, 7, 8, 6, 5, 0], \\ [3, 2, 1, 4, 16, 14, 16, 16, 16, 5], \\ [3, 16, 9, 16, 16, 0, 0, 0, 0, 0], \\ [10, 11, 15, 0, 0, 0, 0, 0, 0, 0] \end{bmatrix}$$

Preprocessing Pipeline Summary

Break the sentence into individual words (or array of words)

"You' ll never know the psychopath sitting next to you."



["You' ll", "never", "know", "the", "psychopath", "sitting", "next", "to", "you"]

In case of multiple sentences, use array of sentences

```
[ [ "you", "to", "next", "sitting", "psychopath", "the", "know", "never", "you'll" ],  
  [ "you", "to", "next", "sitting", "murderer", "the", "know", "never", "you'll" ],  
  [ "you", "to", "next", "sitting", "here", "get", "i", "how'd", "think", "you'll" ],  
    [ "said", "i've", "all", "after", "but" ],  
      [ "forget", "don't", "please" ] ]
```

Assign index/number to each word

('next', 3), ('to', 3), ('you', 3), ('sitting', 3), ("you'll", 3),
('never', 2), ('the', 2), ('know', 2), ('all', 1), ('forget', 1),
("don't", 1), ('psychopath', 1), ('said', 1), ('get', 1), ('please', 1),
("how'd", 1), ('after', 1), ('murderer', 1), ('but', 1),
("i've", 1), ('here', 1), ('i', 1), ('think', 1)

Create a vocabulary (of all the words in the text)

'ZERO',
'next', 'to', 'you', 'sitting', "you'll",
'never', 'the', 'know', 'all', 'forget',
"don't", 'psychopath', 'said', 'get', 'please',
'UNK'

Create a word-index mapping (dictionary)

`"ZERO": 0,`
`"next": 1, "to": 2, "you": 3, "sitting": 4, "you'll": 5,`
`"never": 6, "the": 7, "know": 8, "all": 9, "forget": 10,`
`"don't": 11, "psychopath": 12, "said": 13, "get": 14, "please": 15,`
`"UNK": 16`

Tokenize the sentences

$$\begin{bmatrix} [3, 2, 1, 4, 12, 7, 8, 6, 5], \\ [3, 2, 1, 4, 16, 7, 8, 6, 5], \\ [3, 2, 1, 4, 16, 14, 16, 16, 16, 5], \\ [3, 16, 9, 16, 16], \\ [10, 11, 15] \end{bmatrix}$$

Add zero padding

$$\begin{bmatrix} [3, 2, 1, 4, 12, 7, 8, 6, 5, 0], \\ [3, 2, 1, 4, 16, 7, 8, 6, 5, 0], \\ [3, 2, 1, 4, 16, 14, 16, 16, 16, 5], \\ [3, 16, 9, 16, 16, 0, 0, 0, 0, 0], \\ [10, 11, 15, 0, 0, 0, 0, 0, 0, 0] \end{bmatrix}$$

Convert indices to word vectors

[3, 2, 1, 4, 16, 7, 8, 6, 5, 0]

[illegible]

The embeddings shown above are too simple and we would encourage you to use more complex pre-trained embeddings like GloVe, Word2vec etc. in order to get better results. Figure shows 300d Glove embeddings.

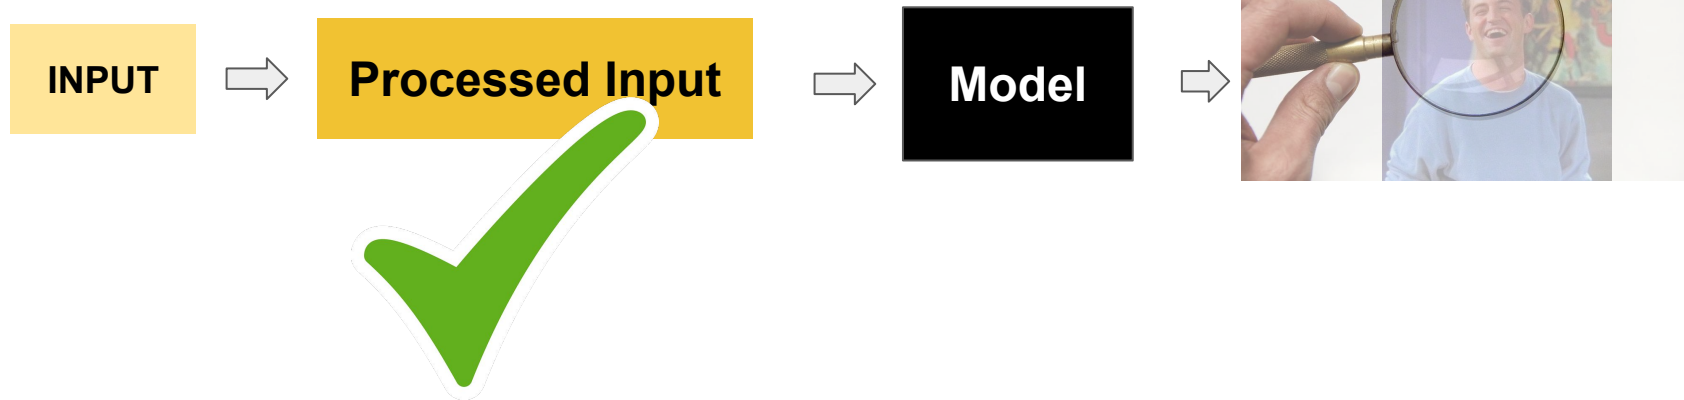
	0	1	2	3	4	5	6	7	8	9	...	290	291	292	
fox	-0.348680	-0.077720	0.177750	-0.094953	-0.452890	0.237790	0.209440	0.037886	0.035064	0.899010	...	-0.283050	0.270240	-0.654800	0.109100
ham	-0.773320	-0.282540	0.580760	0.841480	0.258540	0.585210	-0.021890	-0.463680	0.139070	0.658720	...	0.464470	0.481400	-0.829200	0.354000
brown	-0.374120	-0.076264	0.109260	0.186620	0.029943	0.182700	-0.631980	0.133060	-0.128980	0.603430	...	-0.015404	0.392890	-0.034826	-0.720000
beautiful	0.171200	0.534390	-0.348540	-0.097234	0.101800	-0.170860	0.295650	-0.041816	-0.516550	2.117200	...	-0.285540	0.104670	0.126310	0.120000
jumps	-0.334840	0.215990	-0.350440	-0.260020	0.411070	0.154010	-0.386110	0.206380	0.386700	1.460500	...	-0.107030	-0.279480	-0.186200	-0.540000
eggs	-0.417810	-0.035192	-0.126150	-0.215930	-0.669740	0.513250	-0.797090	-0.068611	0.634660	1.256300	...	-0.232860	-0.139740	-0.681080	-0.370000
beans	-0.423290	-0.264500	0.200870	0.082187	0.066944	1.027600	-0.989140	-0.259950	0.145960	0.766450	...	0.048760	0.351680	-0.786260	-0.360000
sky	0.312550	-0.303080	0.019587	-0.354940	0.100180	-0.141530	-0.514270	0.886110	-0.530540	1.556600	...	-0.667050	0.279110	0.500970	-0.270000
bacon	-0.430730	-0.016025	0.484620	0.101390	-0.299200	0.761820	-0.353130	-0.325290	0.156730	0.873210	...	0.304240	0.413440	-0.540730	-0.030000
breakfast	0.073378	0.227670	0.208420	-0.456790	-0.078219	0.601960	-0.024494	-0.467980	0.054627	2.283700	...	0.647710	0.373820	0.019931	-0.030000
toast	0.130740	-0.193730	0.253270	0.090102	-0.272580	-0.030571	0.096945	-0.115060	0.484000	0.848380	...	0.142080	0.481910	0.045167	0.050000
today	-0.156570	0.594890	-0.031445	-0.077586	0.278630	-0.509210	-0.066350	-0.081890	-0.047986	2.803600	...	-0.326580	-0.413380	0.367910	-0.260000
blue	0.129450	0.036518	0.032298	-0.060034	0.399840	-0.103020	-0.507880	0.076630	-0.422920	0.815730	...	-0.501280	0.169010	0.548250	-0.310000
green	-0.072368	0.233200	0.137260	-0.156630	0.248440	0.349870	-0.241700	-0.091426	-0.530150	1.341300	...	-0.405170	0.243570	0.437300	-0.460000
kings	0.259230	-0.854690	0.360010	-0.642000	0.568530	-0.321420	0.173250	0.133030	-0.089720	1.528600	...	-0.470090	0.063743	-0.545210	-0.190000
dog	-0.057120	0.052685	0.003026	-0.048517	0.007043	0.041856	-0.024704	-0.039783	0.009614	0.308416	...	0.003257	-0.036864	-0.043878	0.000000
sausages	-0.174290	-0.064869	-0.046976	0.287420	-0.128150	0.647630	0.056315	-0.240440	-0.025094	0.502220	...	0.302240	0.195470	-0.653980	-0.290000
lazy	-0.353320	-0.299710	-0.176230	-0.321940	-0.385640	0.586110	0.411160	-0.418680	0.073093	1.486500	...	0.402310	-0.038554	-0.288670	-0.240000
love	0.139490	0.534530	-0.252470	-0.125650	0.048748	0.152440	0.199060	-0.065970	0.128830	2.055900	...	-0.124380	0.178440	-0.099469	0.000000
quick	-0.445630	0.191510	-0.249210	0.465900	0.161950	0.212780	-0.046480	0.021170	0.417660	1.686900	...	-0.329460	0.421860	-0.039543	0.150000

20 rows x 300 columns

Congrats!!!

Now you know how to preprocess text data.

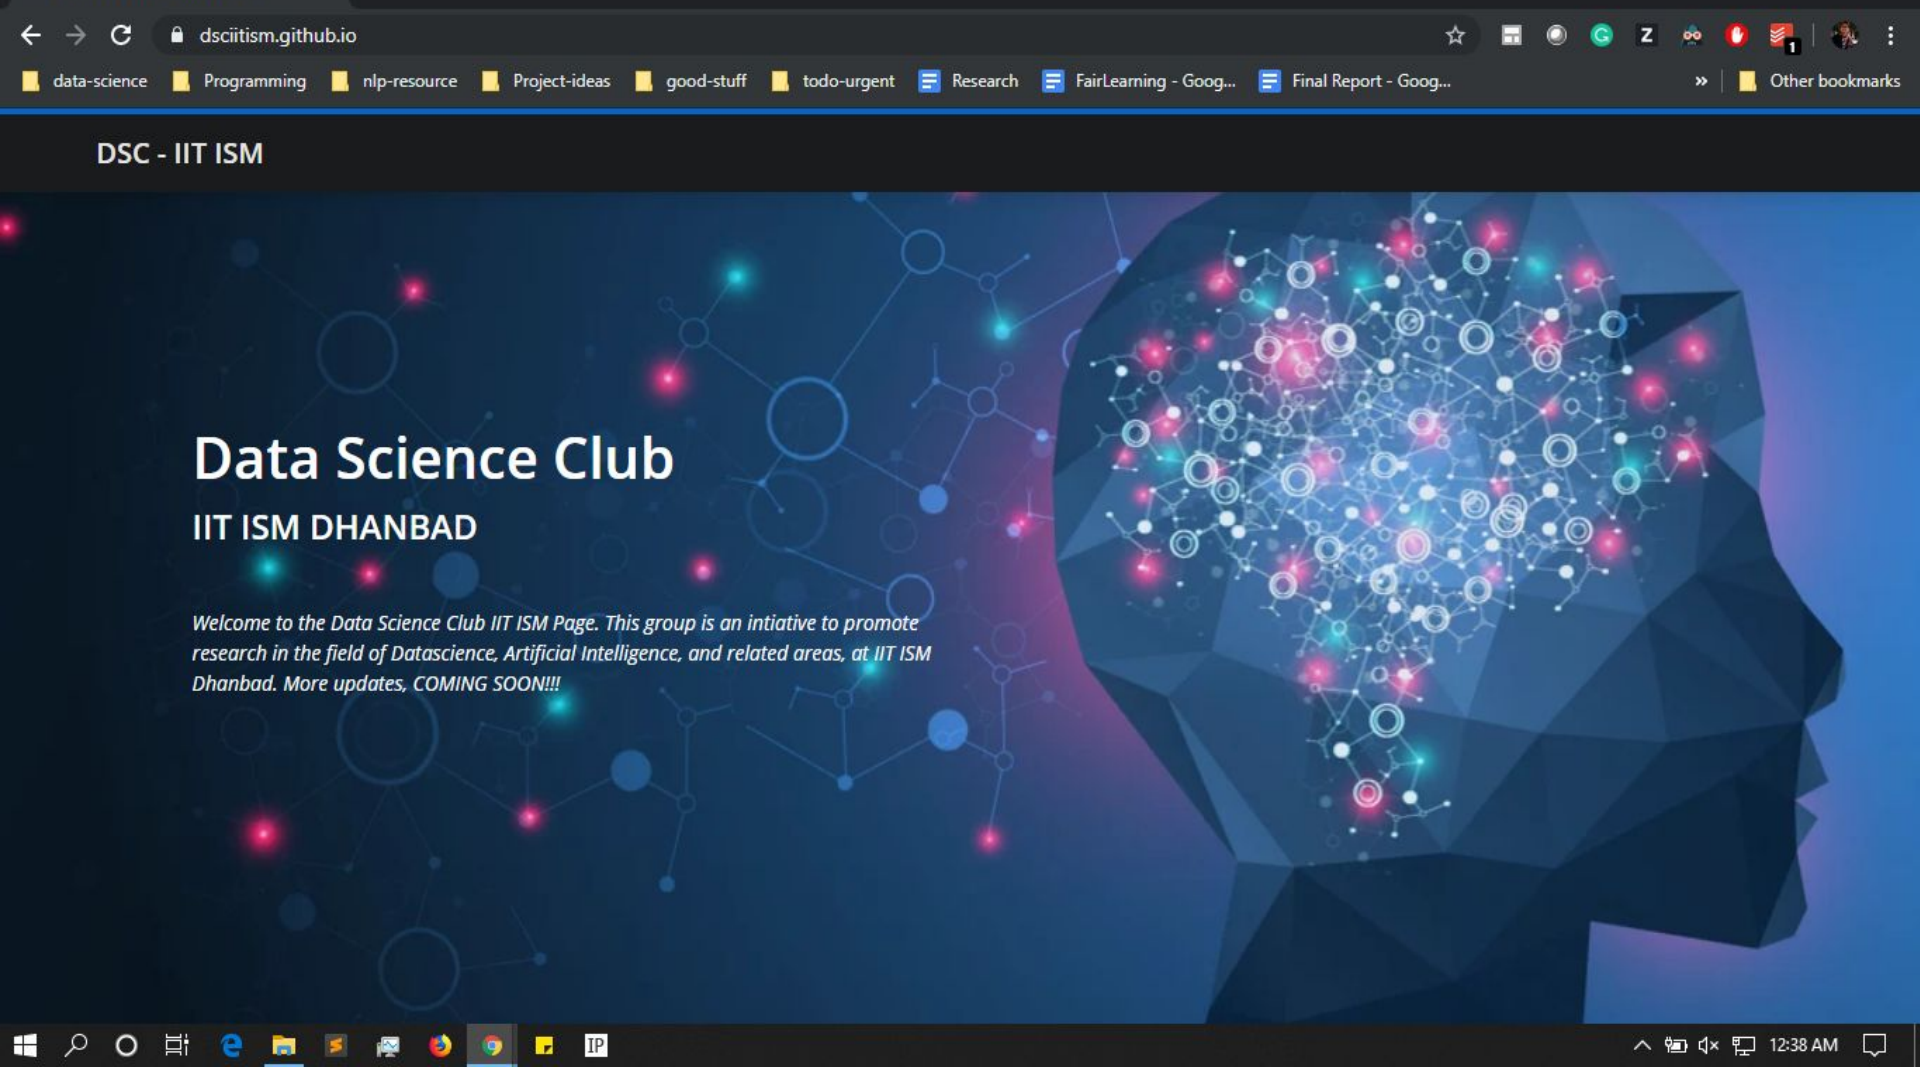
Pipeline:



Datascience Club Page is live !



Kshitij Kumar



DSC - IIT ISM

Data Science Club

IIT ISM DHANBAD

Welcome to the Data Science Club IIT ISM Page. This group is an initiative to promote research in the field of Datascience, Artificial Intelligence, and related areas, at IIT ISM Dhanbad. More updates, COMING SOON!!!!

NLP/ML Workshop

Today's workflow outline:

- Recap
- Different Models
- Baseline Model Tutorial

Text Processing Tutorial

Text changed into numerical form.

Let's address the classifier problem.

Numpy Review

$$X = \begin{bmatrix} x1 \\ x2 \end{bmatrix} \text{ vector (input layer)}$$

$$W = \begin{bmatrix} w1 & w2 \\ w4 & w5 \\ x3 & w6 \end{bmatrix} \text{ matrix (the weights for hidden layer 1)}$$

the output is given by

$$\begin{bmatrix} h1 \\ h2 \\ h3 \end{bmatrix} = \begin{bmatrix} w1 & w2 \\ w4 & w5 \\ x3 & w6 \end{bmatrix} \cdot \begin{bmatrix} x1 \\ x2 \end{bmatrix} \text{ (the product of vector and matrices)}$$


$$\begin{bmatrix} h1 \\ h2 \\ h3 \end{bmatrix} = \begin{bmatrix} h1 \\ h2 \\ h3 \end{bmatrix} + \text{bias}$$

finally,

$$\begin{bmatrix} h1 \\ h2 \\ h3 \end{bmatrix} = f\left(\begin{bmatrix} h1 \\ h2 \\ h3 \end{bmatrix}\right) \text{ (activation step)}$$



Classifier will help us to classify whether or not the given instance belongs to class A or class B.

Note class maybe two (binary), or more (multi class) classification



Different ways to classify

- Draw a straight line
- Draw a curve line
- --




Linear Regression - helps making linear/ polynomial decision boundaries.

Logistic Regression - categorical values

Neural network - helps with more complex boundaries

Linear Regression

Concept



$W x + b$

Loss Function


backpropagation

Loss/Cost Function


$$\text{Least Squared Error} = \frac{1}{2}(h_{\theta}(x^{(i)}) - y^{(i)})^2$$

$$\text{Cost Function} = J(\theta) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2, \quad m = \text{number of sample data}$$

backpropagation



backpropagation




Linear Regression:


Notebook


Logistic Regression


Why Logistic Regression?




Non-linear functions | Need of threshold








Neural Network



Input
(features)

Logistic
classifier

Logistic Regression



Neural Network

Key difference b/w normal data and text

- **SEQUENTIAL ORDER**

House features: (#bedrooms, #windows, #kitchen, #washroom,)

(..., #windows, #bedrooms, #washroom, #kitchen,)

Text feature: (“And”, “I”, “want”, “a”, “million”, “dollar”, “!”)

(“want”, “And”, “million”, “I”, “dollar”, “a”, “!”)


Few other differences

- Long term dependencies
- ...
-


*“Well satisfied with his purchases and feeling very elegant indeed, Babar goes to the photographer to have **his picture** taken.”*

What can help?


RNN (recurrent neural networks) - idea is to use previous input along with current input.



- LSTM (Long Short Term Memory): consists of components (gates) like remember, forget, update , input.



- LSTM (Long Short Term Memory): consists of components (gates) like remember, forget, update , input.




$$z_t = \sigma (W_z \cdot [h_{t-1}, x_t])$$

$$r_t = \sigma (W_r \cdot [h_{t-1}, x_t])$$

$$\tilde{h}_t = \tanh (W \cdot [r_t * h_{t-1}, x_t])$$

$$h_t = (1 - z_t) * h_{t-1} + z_t * \tilde{h}_t$$

Pipeline:



Baseline Tutorial

References:

- <https://distill.pub/2016/handwriting/>
- <https://colah.github.io/posts/2015-08-Understanding-LSTMs/>
- <http://www.wildml.com/2015/09/recurrent-neural-networks-tutorial-part-1-introduction-to-rnns/>
- <https://towardsdatascience.com/natural-language-processing-from-basics-to-using-rnn-and-lstm-ef6779e4ae66>

Courses: [Machine Learning](#) and [deeplearning.ai Specialization](#) by Andrew NG.